

Question 01

```
import math
import matplotlib.pyplot as plt

# Objective function
def revenue(w):
    return 0.5 * w**2 - 30 * w + 100

# Gradient of the objective function
def gradient(w):
    return w - 30

# Hyperparameters
learning_rate = 0.1
momentum = 0.9
iterations = 3

# Initial price
w = 20

# Initialize velocity
v = 0


# Lists to store optimization history for visualization
w_history = [w]
revenue_history = [revenue(w)]

for i in range(iterations):
    v = momentum * v + (1 - momentum) * gradient(w)
    w -= learning_rate * v
    w_history.append(w)
    revenue_history.append(revenue(w))

print("Optimal price:", w)

# Visualization
plt.figure(figsize=(10, 6))
plt.plot(w_history, revenue_history, marker='o', linestyle='-')
plt.title('Optimization Process')
plt.xlabel('Price (w)')
plt.ylabel('Revenue')
plt.grid(True)
plt.show()

Optimal price: 20.55621
```



Question - 2

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

df = pd.read_csv("/content/diabetes.csv")
df.head().T

{"summary":{"\n  \"name\": \"df\",\n  \"rows\": 768,\n  \"fields\": [\n    {\n      \"column\": \"Pregnancies\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 3,\n        \"min\": 0,\n        \"max\": 17,\n        \"num_unique_values\": 17,\n        \"samples\": [\n          6,\n          1,\n          3\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Glucose\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 31,\n        \"min\": 0,\n        \"max\": 199,\n        \"num_unique_values\": 136,\n        \"samples\": [\n          151,\n          101,\n          112\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"BloodPressure\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 19,\n        \"min\": 0,\n        \"max\": 122,\n        \"num_unique_values\": 47,\n
```

```

{"samples": [\n          86,\n          46,\n          85\n        ],\n  "semantic_type": \"\",\n  "description": \"\",\n  "column": \"SkinThickness\",\n  "properties": {\n    \"dtype\": \"number\",\n    \"std\": 15,\n    \"min\": 0,\n    \"max\": 99,\n    \"num_unique_values\": 51,\n    \"samples\": [\n      7,\n      12,\n      48\n    ],\n    \"semantic_type\": \"\",\n    \"description\": \"\",\n    \"column\": \"Insulin\",\n    \"properties\": {\n      \"dtype\": \"number\",\n      \"std\": 115,\n      \"min\": 0,\n      \"max\": 846,\n      \"num_unique_values\": 186,\n      \"samples\": [\n        52,\n        41,\n        183\n      ],\n      \"semantic_type\": \"\",\n      \"description\": \"\",\n      \"column\": \"BMI\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 7.884160320375446,\n        \"min\": 0.0,\n        \"max\": 67.1,\n        \"num_unique_values\": 248,\n        \"samples\": [\n          19.9,\n          31.0,\n          38.1\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\",\n        \"column\": \"DiabetesPedigreeFunction\",\n        \"properties\": {\n          \"dtype\": \"number\",\n          \"std\": 0.3313285950127749,\n          \"min\": 0.078,\n          \"max\": 2.42,\n          \"num_unique_values\": 517,\n          \"samples\": [\n            1.731,\n            0.426,\n            0.138\n          ],\n          \"semantic_type\": \"\",\n          \"description\": \"\",\n          \"column\": \"Age\",\n          \"properties\": {\n            \"dtype\": \"number\",\n            \"std\": 11,\n            \"min\": 21,\n            \"max\": 81,\n            \"num_unique_values\": 52,\n            \"samples\": [\n              60,\n              47,\n              72\n            ],\n            \"semantic_type\": \"\",\n            \"description\": \"\",\n            \"column\": \"Outcome\",\n            \"properties\": {\n              \"dtype\": \"number\",\n              \"std\": 0,\n              \"min\": 0,\n              \"max\": 1,\n              \"num_unique_values\": 2,\n              \"samples\": [\n                0,\n                1\n              ],\n              \"semantic_type\": \"\",\n              \"description\": \"\"\n            }\n          }\n        }\n      ],\n      \"type\": \"dataframe\", \"variable_name\": \"df\"}

```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 768 entries, 0 to 767
```

```
Data columns (total 9 columns):
```

#	Column	Non-Null Count	Dtype
0	Pregnancies	768 non-null	int64
1	Glucose	768 non-null	int64
2	BloodPressure	768 non-null	int64
3	SkinThickness	768 non-null	int64
4	Insulin	768 non-null	int64
5	BMI	768 non-null	float64

6	DiabetesPedigreeFunction	768	non-null	float64
7	Age	768	non-null	int64
8	Outcome	768	non-null	int64

dtypes: float64(2), int64(7)

memory usage: 54.1 KB

df.isna().sum()

Pregnancies	0
Glucose	0
BloodPressure	0
SkinThickness	0
Insulin	0
BMI	0
DiabetesPedigreeFunction	0
Age	0
Outcome	0

dtype: int64

No null values.

df.describe()

```
{
  "summary": {
    "name": "df",
    "rows": 8,
    "fields": [
      {
        "column": "Pregnancies",
        "properties": {
          "dtype": "number",
          "std": 269.85223453356366,
          "min": 0.0,
          "max": 768.0,
          "num_unique_values": 8,
          "samples": [
            3.8450520833333335,
            3.0,
            768.0
          ],
          "semantic_type": "",
          "description": ""
        },
        "column": "Glucose",
        "properties": {
          "dtype": "number",
          "std": 243.73802348295857,
          "min": 0.0,
          "max": 768.0,
          "num_unique_values": 8,
          "samples": [
            120.89453125,
            117.0,
            768.0
          ],
          "semantic_type": "",
          "description": ""
        },
        "column": "BloodPressure",
        "properties": {
          "dtype": "number",
          "std": 252.8525053581062,
          "min": 0.0,
          "max": 768.0,
          "num_unique_values": 8,
          "samples": [
            69.10546875,
            72.0,
            768.0
          ],
          "semantic_type": "",
          "description": ""
        },
        "column": "SkinThickness",
        "properties": {
          "dtype": "number",
          "std": 263.7684730531098,
          "min": 0.0,
          "max": 768.0,
          "num_unique_values": 7,
          "samples": [
            768.0,
            20.536458333333332,
            32.0
          ],
          "semantic_type": "",
          "description": ""
        },
        "column": "Insulin",
        "properties": {
          "dtype": "number",
          "std": 350.26059167945886,
          "min": 0.0,
          "max": 768.0,
          "num_unique_values": 7,
          "samples": [
            768.0,
            20.536458333333332,
            32.0
          ],
          "semantic_type": "",
          "description": ""
        }
      ]
    }
  }
}
```

```

{"max": 846.0, "num_unique_values": 7, "samples": [768.0, 79.79947916666667, 127.25], "semantic_type": "", "description": "", "column": "BMI", "properties": {"dtype": "number", "std": 262.05117817552093, "min": 0.0, "max": 768.0, "num_unique_values": 8, "samples": [31.992578124999998, 32.0, 768.0]}, "semantic_type": "", "description": "", "column": "DiabetesPedigreeFunction", "properties": {"dtype": "number", "std": 271.3005221658502, "min": 0.078, "max": 768.0, "num_unique_values": 8, "samples": [0.47187630208333325, 0.3725, 768.0]}, "semantic_type": "", "description": "", "column": "Age", "properties": {"dtype": "number", "std": 260.1941178528413, "min": 11.760231540678685, "max": 768.0, "num_unique_values": 8, "samples": [33.240885416666664, 29.0, 768.0]}, "semantic_type": "", "description": "", "column": "Outcome", "properties": {"dtype": "number", "std": 271.3865920388932, "min": 0.0, "max": 768.0, "num_unique_values": 5, "samples": [0.3489583333333333, 1.0, 0.47695137724279896]}, "semantic_type": "", "description": ""}, "type": "dataframe"}

```

df.shape

(768, 9)

Numerical features

df.describe(exclude=['O'])

```

{"summary": {"name": "df", "rows": 8, "fields": [{"column": "Pregnancies", "properties": {"dtype": "number", "std": 269.85223453356366, "min": 0.0, "max": 768.0, "num_unique_values": 8, "samples": [3.8450520833333335, 3.0, 768.0]}, "semantic_type": "", "description": "", "column": "Glucose", "properties": {"dtype": "number", "std": 243.73802348295857, "min": 0.0, "max": 768.0, "num_unique_values": 8, "samples": [120.89453125, 117.0, 768.0]}, "semantic_type": "", "description": "", "column":

```

```

\"BloodPressure\", \n      \"properties\": { \n      \"dtype\":
\"number\", \n      \"std\": 252.8525053581062, \n      \"min\":
0.0, \n      \"max\": 768.0, \n      \"num_unique_values\": 8, \n
\"samples\": [ \n      69.10546875, \n      72.0, \n
768.0 \n      ], \n      \"semantic_type\": \"\", \n
\"description\": \"\" \n      } \n      }, \n      { \n      \"column\":
\"SkinThickness\", \n      \"properties\": { \n      \"dtype\":
\"number\", \n      \"std\": 263.7684730531098, \n      \"min\":
0.0, \n      \"max\": 768.0, \n      \"num_unique_values\": 7, \n
\"samples\": [ \n      768.0, \n      20.536458333333332, \n
32.0 \n      ], \n      \"semantic_type\": \"\", \n
\"description\": \"\" \n      } \n      }, \n      { \n      \"column\":
\"Insulin\", \n      \"properties\": { \n      \"dtype\": \"number\", \n
      \"std\": 350.26059167945886, \n      \"min\": 0.0, \n
      \"max\": 846.0, \n      \"num_unique_values\": 7, \n
\"samples\": [ \n      768.0, \n      79.79947916666667, \n
127.25 \n      ], \n      \"semantic_type\": \"\", \n
\"description\": \"\" \n      } \n      }, \n      { \n      \"column\":
\"BMI\", \n      \"properties\": { \n      \"dtype\": \"number\", \n
      \"std\": 262.05117817552093, \n      \"min\": 0.0, \n
      \"max\": 768.0, \n      \"num_unique_values\": 8, \n
\"samples\": [ \n      31.992578124999998, \n      32.0, \n
768.0 \n      ], \n      \"semantic_type\": \"\", \n
\"description\": \"\" \n      } \n      }, \n      { \n      \"column\":
\"DiabetesPedigreeFunction\", \n      \"properties\": { \n      \"dtype\": \"number\", \n
      \"std\": 271.3005221658502, \n      \"min\": 0.078, \n
      \"max\": 768.0, \n      \"num_unique_values\": 8, \n
\"samples\": [ \n      0.47187630208333325, \n      0.3725, \n
768.0 \n      ], \n      \"semantic_type\": \"\", \n
\"description\": \"\" \n      } \n      }, \n      { \n      \"column\":
\"Age\", \n      \"properties\": { \n      \"dtype\": \"number\", \n
      \"std\": 260.1941178528413, \n      \"min\": 11.760231540678685, \n
      \"max\": 768.0, \n      \"num_unique_values\": 8, \n
\"samples\": [ \n      33.240885416666664, \n      29.0, \n
768.0 \n      ], \n      \"semantic_type\": \"\", \n
\"description\": \"\" \n      } \n      }, \n      { \n      \"column\":
\"Outcome\", \n      \"properties\": { \n      \"dtype\": \"number\", \n
      \"std\": 271.3865920388932, \n      \"min\": 0.0, \n
      \"max\": 768.0, \n      \"num_unique_values\": 5, \n
\"samples\": [ \n      0.3489583333333333, \n      1.0, \n
0.47695137724279896 \n      ], \n      \"semantic_type\": \"\", \n
\"description\": \"\" \n      } \n      } \n      ], \n      \"type\": \"dataframe\"}

```

No categorical columns, so no need for hot encoding

```

#lib for model building & eval
import tensorflow as tf
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, precision_score,

```

```

recall_score, f1_score, confusion_matrix, roc_curve, auc
import matplotlib.pyplot as plt

# split data into features and target
X = df.drop('Outcome', axis=1)
y = df['Outcome']

# split data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

```

Base Model

```

# neural network model
model = tf.keras.models.Sequential([
    tf.keras.layers.Dense(64, activation='relu',
input_shape=(X_train.shape[1],)),
    tf.keras.layers.Dense(32, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])

# model.compile
model.compile(optimizer='adam',
              loss='binary_crossentropy',
              metrics=['accuracy'])

# model train
hist = model.fit(X_train, y_train, epochs=50, batch_size=32,
validation_split=0.2, verbose=0)

# model eval
y_pred = model.predict(X_test)
y_pred = (y_pred > 0.5).astype(int)

acc = accuracy_score(y_test, y_pred)
prec = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print("Accuracy:", acc)
print("Precision:", prec)
print("Recall:", recall)
print("F1 Score:", f1)

5/5 [=====] - 0s 4ms/step
Accuracy: 0.6298701298701299
Precision: 0.4722222222222222
Recall: 0.3090909090909091
F1 Score: 0.37362637362637363


```

```
# Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:")
print(cm)

Confusion Matrix:
[[80 19]
 [38 17]]

# ROC Curve
fpr, tpr, thresholds = roc_curve(y_test, y_pred)
roc_auc = auc(fpr, tpr)


plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC')
plt.legend(loc="lower right")
plt.show()
```

```
# Learning Curves
plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
plt.plot(hist.history['accuracy'], label='accuracy')
plt.plot(hist.history['val_accuracy'], label='val_accuracy')
plt.title('Training and Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(hist.history['loss'], label='loss')
plt.plot(hist.history['val_loss'], label='val_loss')
plt.title('Training and Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()

plt.show()
```



User defined code for evaluation, Conf matrix, ROC curve, Learning curve for reusability.

```
# Evaluate the performance of models
def evaluate_model(model, X_test, y_test):
    y_pred = model.predict(X_test)
    y_pred = (y_pred > 0.5).astype(int)

    accuracy = accuracy_score(y_test, y_pred)
    precision = precision_score(y_test, y_pred)
    recall = recall_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred)

    return accuracy, precision, recall, f1

# Confusion Matrix
def plot_confusion_matrix(model, X_test, y_test):
    y_pred = model.predict(X_test)
    y_pred = (y_pred > 0.5).astype(int)
    cm = confusion_matrix(y_test, y_pred)
    plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Blues)
    plt.title('Confusion Matrix')
    plt.colorbar()
    plt.xlabel('Predicted')
    plt.ylabel('Actual')
    plt.xticks([0, 1], ['Negative', 'Positive'])
    plt.yticks([0, 1], ['Negative', 'Positive'])

    thresh = cm.max() / 2.
    for i in range(cm.shape[0]):
        for j in range(cm.shape[1]):
```

```

        plt.text(j, i, format(cm[i, j], 'd'),
horizontalalignment="center", color="white" if cm[i, j] > thresh else
"black")

plt.show()

# ROC Curve
def plot_roc_curve(model, X_test, y_test):
    y_pred = model.predict(X_test)
    fpr, tpr, thresholds = roc_curve(y_test, y_pred)
    roc_auc = auc(fpr, tpr)
    plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve
(area = %0.2f)' % roc_auc)
    plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.05])
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('ROC')
    plt.legend(loc="lower right")
    plt.show()

# Learning Curves
def plot_learning_curves(history, title):
    plt.plot(history.history['accuracy'], label='accuracy')
    plt.plot(history.history['val_accuracy'], label='val_accuracy')
    plt.title(title)
    plt.xlabel('Epoch')
    plt.ylabel('Accuracy')
    plt.legend()
    plt.show()

```

L1 Model

```

# L1 regularization
model_l1 = tf.keras.models.Sequential([
    tf.keras.layers.Dense(64, activation='relu',
kernel_regularizer=tf.keras.regularizers.l1(0.01),
input_shape=(X_train.shape[1],)),
    tf.keras.layers.Dense(32, activation='relu',
kernel_regularizer=tf.keras.regularizers.l1(0.01)),
    tf.keras.layers.Dense(1, activation='sigmoid')
])

# model
model_l1.compile(optimizer='adam',
                 loss='binary_crossentropy',
                 metrics=['accuracy'])

```

```

# Train
hist_l1 = model_l1.fit(X_train, y_train, epochs=50, batch_size=32,
validation_split=0.2, verbose=0)


#eval
accuracy_l1, precision_l1, recall_l1, f1_l1 = evaluate_model(model_l1,
X_test, y_test)
print("L1 Regularization:")
print("Accuracy:", accuracy_l1)
print("Precision:", precision_l1)
print("Recall:", recall_l1)
print("F1 Score:", f1_l1)

5/5 [=====] - 0s 4ms/step
L1 Regularization:
Accuracy: 0.6688311688311688
Precision: 0.5416666666666666
Recall: 0.4727272727272727
F1 Score: 0.5048543689320388

print("\nConfusion Matrix - L1 Regularization:")
plot_confusion_matrix(model_l1, X_test, y_test)


Confusion Matrix - L1 Regularization:
5/5 [=====] - 0s 2ms/step

```



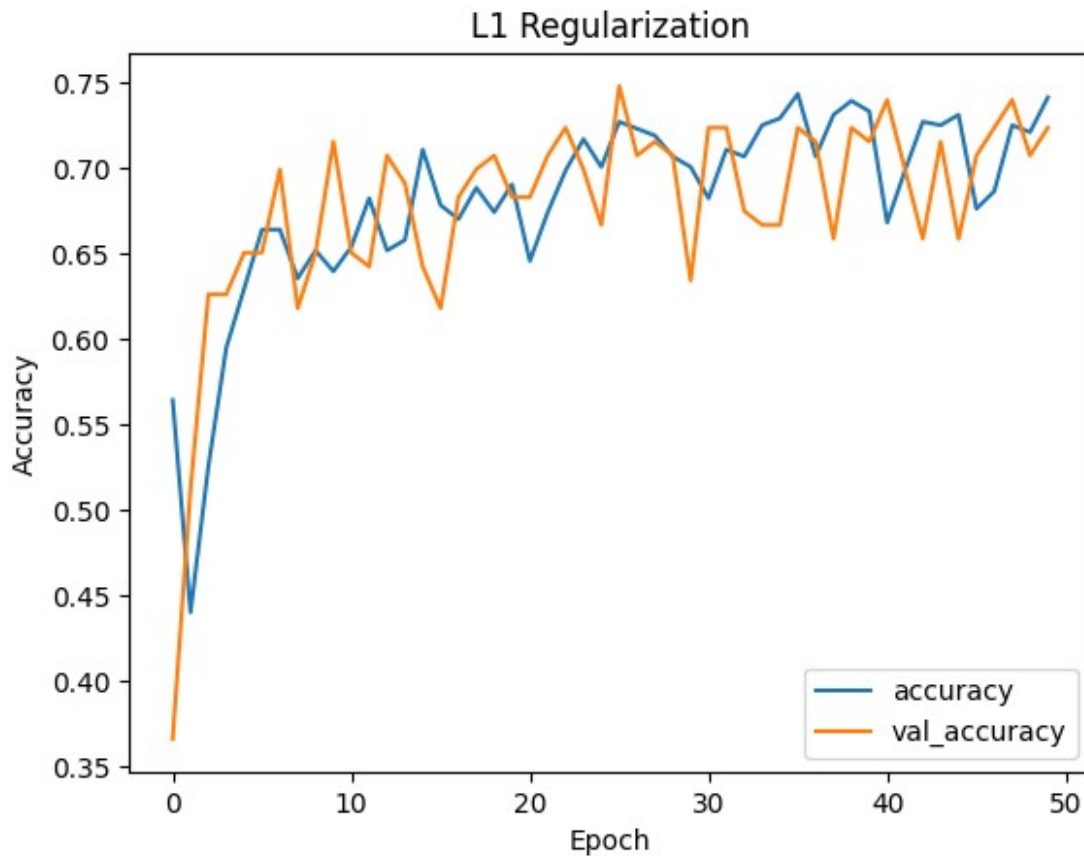
```
print("\nROC Curve - L1 Regularization:")  
plot_roc_curve(model_l1, X_test, y_test)
```

```
ROC Curve - L1 Regularization:  
5/5 [=====] - 0s 3ms/step
```



```
print("\nLearning Curves - L1 Regularization:")  
plot_learning_curves(hist_l1, "L1 Regularization")
```

Learning Curves - L1 Regularization:



L2

```
# L2 regularization
model_l2 = tf.keras.models.Sequential([
    tf.keras.layers.Dense(64, activation='relu',
        kernel_regularizer=tf.keras.regularizers.l2(0.01),
        input_shape=(X_train.shape[1],)),
    tf.keras.layers.Dense(32, activation='relu',
        kernel_regularizer=tf.keras.regularizers.l2(0.01)),
    tf.keras.layers.Dense(1, activation='sigmoid')
])

# model
model_l2.compile(optimizer='adam',
                 loss='binary_crossentropy',
                 metrics=['accuracy'])

# Train
hist_l2 = model_l2.fit(X_train, y_train, epochs=50, batch_size=32,
                      validation_split=0.2, verbose=0)

#eval
accuracy_l2, precision_l2, recall_l2, f1_l2 = evaluate_model(model_l2,
X_test, y_test)
```

```

print("\nL2 Regularization:")
print("Accuracy:", accuracy_l2)
print("Precision:", precision_l2)
print("Recall:", recall_l2)
print("F1 Score:", f1_l2)


5/5 [=====] - 0s 3ms/step

L2 Regularization:
Accuracy: 0.6623376623376623
Precision: 0.6363636363636364
Recall: 0.12727272727272726
F1 Score: 0.2121212121212121

print("\nConfusion Matrix - L2 Regularization:")
plot_confusion_matrix(model_l2, X_test, y_test)

Confusion Matrix - L2 Regularization:
5/5 [=====] - 0s 3ms/step

```




```

print("\nROC Curve - L2 Regularization:")
plot_roc_curve(model_l2, X_test, y_test)

```



ROC Curve - L2 Regularization:

5/5 [=====] - 0s 6ms/step



```
print("\nLearning Curves - L2 Regularization:")  
plot_learning_curves(hist_l2, "L2 Regularization")
```

Learning Curves - L2 Regularization:



Elastic net

```
# Elastic Net regularization
model_elastic_net = tf.keras.models.Sequential([
    tf.keras.layers.Dense(64, activation='relu',
        kernel_regularizer=tf.keras.regularizers.l1_l2(l1=0.01, l2=0.01),
        input_shape=(X_train.shape[1],)),
    tf.keras.layers.Dense(32, activation='relu',
        kernel_regularizer=tf.keras.regularizers.l1_l2(l1=0.01, l2=0.01)),
    tf.keras.layers.Dense(1, activation='sigmoid')
])

# model
model_elastic_net.compile(optimizer='adam',
    loss='binary_crossentropy',
    metrics=['accuracy'])

# Train the model
hist_elastic_net = model_elastic_net.fit(X_train, y_train, epochs=50,
    batch_size=32, validation_split=0.2, verbose=0)

#eval
accuracy_elastic_net, precision_elastic_net, recall_elastic_net,
```


```
f1_elastic_net = evaluate_model(model_elastic_net, X_test, y_test)
print("\nElastic Net Regularization:")
print("Accuracy:", accuracy_elastic_net)
print("Precision:", precision_elastic_net)
print("Recall:", recall_elastic_net)
print("F1 Score:", f1_elastic_net)
```

5/5 [=====] - 0s 3ms/step

Elastic Net Regularization:
Accuracy: 0.6753246753246753
Precision: 0.5316455696202531
Recall: 0.7636363636363637
F1 Score: 0.6268656716417911


```
print("\nConfusion Matrix - Elastic Net Regularization:")
plot_confusion_matrix(model_elastic_net, X_test, y_test)
```

Confusion Matrix - Elastic Net Regularization:
5/5 [=====] - 0s 3ms/step



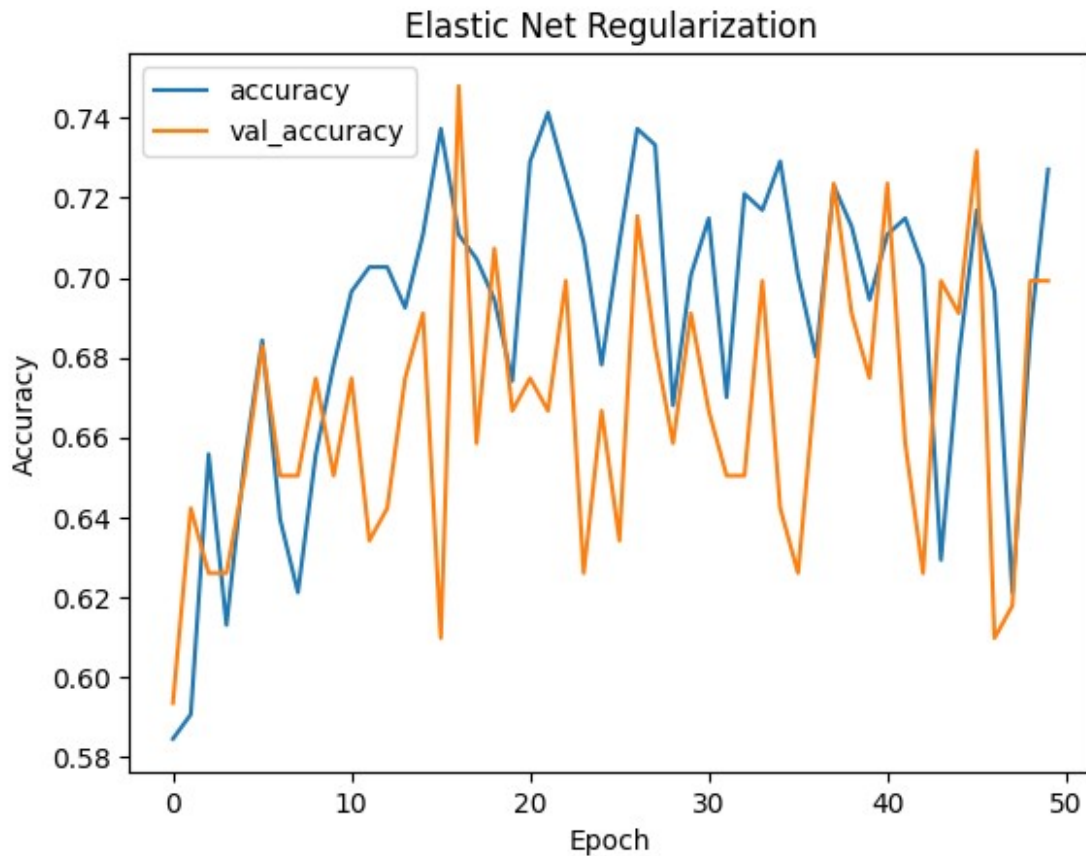
```
print("\nROC Curve - Elastic Net Regularization:")
plot_roc_curve(model_elastic_net, X_test, y_test)
```

ROC Curve - Elastic Net Regularization:
5/5 [=====] - 0s 3ms/step



```
print("\nLearning Curves - Elastic Net Regularization:")
plot_learning_curves(hist_elastic_net, "Elastic Net Regularization")
```

Learning Curves - Elastic Net Regularization:



Comparison

```
# Comparison of metrics
labels = ['Accuracy', 'Precision', 'Recall', 'F1 Score']
metrics_nn = [accuracy_score(y_test, y_pred), precision_score(y_test,
y_pred), recall_score(y_test, y_pred), f1_score(y_test, y_pred)]
metrics_l1 = [accuracy_l1, precision_l1, recall_l1, f1_l1]
metrics_l2 = [accuracy_l2, precision_l2, recall_l2, f1_l2]
metrics_elastic_net = [accuracy_elastic_net, precision_elastic_net,
recall_elastic_net, f1_elastic_net]

x = np.arange(len(labels))
width = 0.2

fig, ax = plt.subplots(figsize=(10, 6))

rects1 = ax.bar(x - width, metrics_nn, width, label='NN')
rects2 = ax.bar(x, metrics_l1, width, label='L1')
rects3 = ax.bar(x + width, metrics_l2, width, label='L2')
rects4 = ax.bar(x + 2*width, metrics_elastic_net, width,
label='Elastic Net')


ax.set_ylabel('Scores')
```

```

ax.set_title('Comparison of Metrics')
ax.set_xticks(x)
ax.set_xticklabels(labels)
ax.legend()

plt.show()

```



```

# Training loss comparison
plt.figure(figsize=(10, 6))
plt.plot(hist.history['loss'], label='NN Loss', linestyle='--')
plt.plot(hist_l1.history['loss'], label='L1 Loss', linestyle='--')
plt.plot(hist_l2.history['loss'], label='L2 Loss', linestyle='--')
plt.plot(hist_elastic_net.history['loss'], label='Elastic Net Loss',
linestyle=':')
plt.title('Training Loss Comparison')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.grid(True)
plt.show()

```

